


ORNL Brand Agent Model Cost Comparison

ORNL Brand Agent V2.3.2 | Heavy-duty EV graphic request
GPT-5.6 Sol settings | Same project brief | One measured run per setting | 500 credits per user per month

Per-run price spans 5x: \$0.40 to \$2.00. Monthly capacity falls from 50 to 10 comparable runs.

Displayed estimate per run - lowest to highest



Bar length shows the dashboard's displayed dollar estimate. The right column shows floor(500 credits / credits per run).

Cost, capacity, and quality snapshot

SETTING	CREDITS	EST. \$	ALLOWANCE	MAX RUNS	SINGLE-RUN QA
Instant	10	\$0.40	2%	50	Not scored
Thinking - Mini	10	\$0.40	2%	50	Rejected / 0 passed
Light	20	\$0.80	4%	25	Concept / 0 passed
Thinking - Standard	20	\$0.80	4%	25	Concept / 0 passed
Pro + Pro thinking	50	\$2.00	10%	10	Rejected / 0 passed

Leadership summary

Lowest cost: Instant and Thinking-Mini cost \$0.40 per run and support 50 monthly runs. Instant remains unscored; Mini failed the quality gate.	No premium payoff yet: Light, Standard, and Pro all recorded zero passing outputs. The \$2.00 Pro run did not show a usable-quality advantage.	Decision: Repeat each setting against one fixed QA rubric, then compare credits per passing output before selecting a default.
---	--	--

Method: after-minus-before dashboard deltas; n=1 per setting. Four tests were scored and none passed; Instant was not scored. Quality evidence sources differ, so these results are directional rather than statistically predictive. Displayed dollar estimates may not equal incremental invoice charges.